

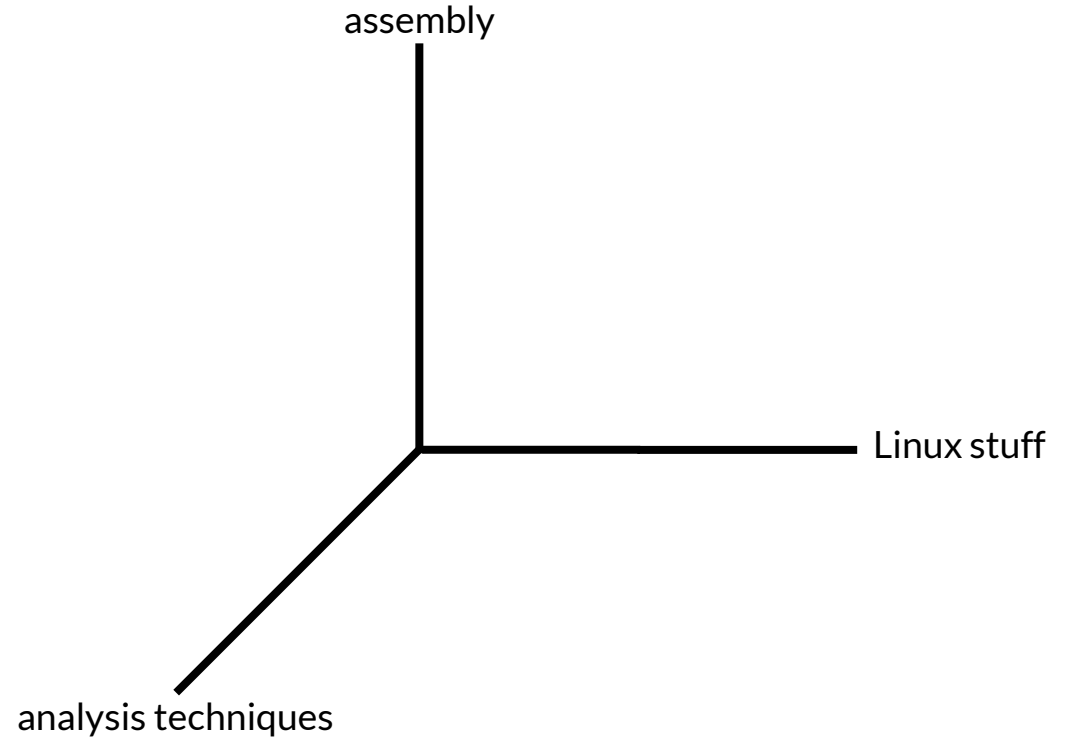
CSC 6580

Spring 2020

Instructor: Stacy Prowell



The Class in 3D



Assembly: Jump Tables



Case Statements

```
switch(x) {  
    case 0: f();  
    case 1: g(); break;  
    case 2: h(); break;  
    default: j();  
}
```

What will the compiler do here?

It will *probably* just create a sequence of branches as this is very simple code, but it might also create a *jump table*.



Case Statements and Jump Tables

```
switch(x) {  
    case 0: f();  
    case 1: g(); break;  
    case 2: h(); break;  
    default: j();  
}
```

We can implement this as a **jump table**. This is a common structure.

```
section .data  
jumptab: dq targetf, targetg, targeth  
  
section .text  
    ; ...  
    cmp rax, 3  
    ja .def  
    jmp [jumptab + rax*8]  
  
.def:  
    jmp targetj
```

ELF, the entry point, and main



Very Simple File

Let's examine a very simple program: hello world.

We can compile this with:

```
ac () { eval $( head -1 $1 | cut -c3- ) }
```

This produces a file about 9,000 bytes in length.

```
; nasm -f elf64 hello.asm && ld -o hello hello.o
```

```
section .text  
global _start
```

```
_start:
```

```
mov rdi, 1  
mov rsi, msg1  
mov rdx, msg1.len  
mov rax, 1  
syscall
```

```
mov rdi, 1  
mov rsi, msg2  
mov rdx, msg2.len  
mov rax, 1  
syscall
```

```
mov rdi, 0  
mov rax, 60  
syscall  
hlt
```

```
msg1:      section .data nowrite align=16  
.len:      db 'Hello world!',10  
           equ $-msg1  
msg2:      db 'Goodbye world!',10  
.len:      equ $-msg2
```



Very Simple File

Define a value. The dollar sign (\$) refers to the *current address*.

```
; nasm -f elf64 hello.asm && ld -o hello hello.o
```

```
section .text  
global _start
```

```
_start:
```

```
mov rdi, 1  
mov rsi, msg1  
mov rdx, msg1.len  
mov rax, 1  
syscall
```

```
mov rdi, 1  
mov rsi, msg2  
mov rdx, msg2.len  
mov rax, 1  
syscall
```

```
mov rdi, 0  
mov rax, 60  
syscall  
hlt
```

```
section .data nowrite align=16  
msg1: db 'Hello world!',10  
.len: equ $-msg1  
msg2: db 'Goodbye world!',10  
.len: equ $-msg2
```




Very Simple File

You can specify additional properties for your sections. Here we make `.data` read-only (it is normally writeable) and force it to be aligned on a 16-byte boundary.

```
; nasm -f elf64 hello.asm && ld -o hello hello.o
```

```
section .text  
global _start
```

```
_start:
```

```
mov rdi, 1  
mov rsi, msg1  
mov rdx, msg1.len  
mov rax, 1  
syscall
```

```
mov rdi, 1  
mov rsi, msg2  
mov rdx, msg2.len  
mov rax, 1  
syscall
```

```
mov rdi, 0  
mov rax, 60  
syscall  
hlt
```

```
msg1:  
.len: equ $-msg1  
msg2:  
.len: equ $-msg2  
  
section .data nowrite align=16  
db 'Hello world!',10  
db 'Goodbye world!',10
```



Very Simple File

We can get an idea of what's in it with `objdump`.

```
$ objdump -s hello
```

```
hello:      file format elf64-x86-64
```

```
Contents of section .text:
```

```
401000 bf010000 0048be00 20400000 000000ba  ....H.. @.....
401010 0d000000 b8010000 000f05bf 01000000  .....
401020 48be0d20 40000000 0000ba0f 000000b8  H.. @.....
401030 01000000 0f05bf00 000000b8 3c000000  .....<...
401040 0f05f4                                ...
```

```
Contents of section .data:
```

```
402000 48656c6c 6f20776f 726c6421 0a476f6f  Hello world!.Goo
402010 64627965 20776f72 6c64210a                dbye world!.
```



Very Simple File

We can get an idea of what's in it with `objdump`.

It's *really simple*. There are just two sections: `.text` and `.data`.

Sections eventually get mapped to *segments*. You can have any sections you want, and can specify permissions on them.

```
section .special write align=4
```

```
$ objdump -s hello
```

```
hello:      file format elf64-x86-64
```

```
Contents of section .text:
```

```
401000 bf010000 0048be00 20400000 000000ba .....H.. @.....
401010 0d000000 b8010000 000f05bf 01000000 .....
401020 48be0d20 40000000 0000ba0f 000000b8 H.. @.....
401030 01000000 0f05bf00 000000b8 3c000000 .....<...
401040 0f05f4                                     ...
```

```
Contents of section .data:
```

```
402000 48656c6c 6f20776f 726c6421 0a476f6f Hello world!.Goo
402010 64627965 20776f72 6c64210a      dbye world!.
```



Very Simple File

We can get an idea of what's in it with `objdump`.

It's *really simple*. There are just two sections: `.text` and `.data`.

The numbers at the start of each line are *virtual addresses* and *not* offsets into the file.

```
$ objdump -s hello
```

```
hello:      file format elf64-x86-64
```

```
Contents of section .text:
```

```
401000 bf010000 0048be00 20400000 000000ba .....H.. @.....
401010 0d000000 b8010000 000f05bf 01000000 .....
401020 48be0d20 40000000 0000ba0f 000000b8 H.. @.....
401030 01000000 0f05bf00 000000b8 3c000000 .....<...
401040 0f05f4                                     ...
```

```
Contents of section .data:
```

```
402000 48656c6c 6f20776f 726c6421 0a476f6f Hello world!.Goo
402010 64627965 20776f72 6c64210a      dbye world!.
```



Very Simple File

We can get an idea of what's in it with `objdump`.

It's *really simple*. There are just two sections: `.text` and `.data`.

The numbers at the start of each line are *virtual addresses* and *not* offsets into the file.

You can use `-F` to see the file offsets.

```
$ objdump -s hello -F
```

```
hello:      file format elf64-x86-64
```

```
Contents of section .text: (Starting at file offset: 0x1000)
```

```
401000 bf010000 0048be00 20400000 000000ba  ....H.. @.....
401010 0d000000 b8010000 000f05bf 01000000  .....
401020 48be0d20 40000000 0000ba0f 000000b8  H.. @.....
401030 01000000 0f05bf00 000000b8 3c000000  .....<...
401040 0f05f4                                ...
```

```
Contents of section .data: (Starting at file offset: 0x2000)
```

```
402000 48656c6c 6f20776f 726c6421 0a476f6f  Hello world!.Goo
402010 64627965 20776f72 6c64210a                dbye world!.
```



Very Simple File

We can get an idea of what's in it with `objdump`.

It's *really simple*. There are just two sections: `.text` and `.data`.

The numbers at the start of each line are *virtual addresses* and *not* offsets into the file.

The `-s` flag displays all *non-empty* sections. If you just want a few, use `-j` to specify those sections.

```
$ objdump -s hello -j .text
```

```
hello:      file format elf64-x86-64
```

```
Contents of section .text:
```

```
401000 bf010000 0048be00 20400000 000000ba .....H.. @.....
401010 0d000000 b8010000 000f05bf 01000000 .....
401020 48be0d20 40000000 0000ba0f 000000b8 H.. @.....
401030 01000000 0f05bf00 000000b8 3c000000 .....<...
401040 0f05f4    ...
```



Very Simple File

We can get the file headers with `objdump -f`, or (better yet) with `readelf -h`.

There are actually 6 sections in this file!

```
$ readelf -h hello
```

```
ELF Header:
  Magic:   7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00
  Class:                               ELF64
  Data:                               2's complement, little
  endian
  Version:                             1 (current)
  OS/ABI:                               UNIX - System V
  ABI Version:                           0
  Type:                                EXEC (Executable file)
  Machine:                             Advanced Micro Devices
  X86-64
  Version:                             0x1
  Entry point address:                  0x401000
  Start of program headers:             64 (bytes into file)
  Start of section headers:            8616 (bytes into file)
  Flags:                                0x0
  Size of this header:                  64 (bytes)
  Size of program headers:              56 (bytes)
  Number of program headers:             3
  Size of section headers:              64 (bytes)
  Number of section headers:            6
  Section header string table index: 5
```



Very Simple File

Here are the six sections, as per `readelf`.

```
$ readelf -S hello
```

There are 6 section headers, starting at offset 0x21a8:

Section Headers:

[Nr]	Name	Type	Address	Offset
	Size	EntSize	Flags Link Info Align	
[0]		NULL	0000000000000000	00000000
	0000000000000000	0000000000000000	0 0	0
[1]	.text	PROGBITS	0000000000401000	00001000
	0000000000000043	0000000000000000	AX 0 0	16
[2]	.data	PROGBITS	0000000000402000	00002000
	000000000000001c	0000000000000000	A 0 0	16
[3]	.symtab	SYMTAB	0000000000000000	00002020
	00000000000000120	0000000000000018	4 8	8
[4]	.strtab	STRTAB	0000000000000000	00002140
	000000000000003f	0000000000000000	0 0	1
[5]	.shstrtab	STRTAB	0000000000000000	0000217f
	0000000000000027	0000000000000000	0 0	1

Key to Flags:

W (write), A (alloc), X (execute), M (merge), S (strings), I (info),
L (link order), O (extra OS processing required), G (group), T
(TLS),
C (compressed), x (unknown), o (OS specific), E (exclude),
l (large), p (processor specific)



Very Simple File

The entry point is the (virtual) address where the program begins executing. In our case it is `0x401000`.

```
$ readelf -h hello
```

```
ELF Header:
  Magic:   7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00
  Class:                               ELF64
  Data:                               2's complement, little
  endian
  Version:                             1 (current)
  OS/ABI:                               UNIX - System V
  ABI Version:                           0
  Type:                                 EXEC (Executable file)
  Machine:                               Advanced Micro Devices
  X86-64
  Version:                               0x1
  Entry point address:                   0x401000
  Start of program headers:              64 (bytes into file)
  Start of section headers:              8616 (bytes into file)
  Flags:                                 0x0
  Size of this header:                    64 (bytes)
  Size of program headers:                56 (bytes)
  Number of program headers:              3
  Size of section headers:                64 (bytes)
  Number of section headers:              6
  Section header string table index:      5
```



Very Simple File

The entry point is the (virtual) address where the program begins executing. In our case it is `0x401000`.

Dumping the symbol table, we see that this is the address of `_start`.

```
$ readelf -s hello
```

Symbol table '.symtab' contains 12 entries:

Num:	Value	Size	Type	Bind	Vis	Ndx	Name
0:	0000000000000000	0	NOTYPE	LOCAL	DEFAULT	UND	
1:	0000000000401000	0	SECTION	LOCAL	DEFAULT	1	
2:	0000000000402000	0	SECTION	LOCAL	DEFAULT	2	
3:	0000000000000000	0	FILE	LOCAL	DEFAULT	ABS	hello.asm
4:	0000000000402000	0	NOTYPE	LOCAL	DEFAULT	2	msg1
5:	000000000000000d	0	NOTYPE	LOCAL	DEFAULT	ABS	msg1.len
6:	000000000040200d	0	NOTYPE	LOCAL	DEFAULT	2	msg2
7:	000000000000000f	0	NOTYPE	LOCAL	DEFAULT	ABS	msg2.len
8:	0000000000401000	0	NOTYPE	GLOBAL	DEFAULT	1	_start
9:	000000000040201c	0	NOTYPE	GLOBAL	DEFAULT	2	__bss_start
10:	000000000040201c	0	NOTYPE	GLOBAL	DEFAULT	2	_edata
11:	0000000000402020	0	NOTYPE	GLOBAL	DEFAULT	2	_end



Find the Entry Point

Let's try this for a real file: `ls`.

First... there is no `_start` symbol!
(You can do this, too, with a linker script.)

There must still be an entry point. Let's find it.

```
$ readelf -s `which ls` | grep _start
53: 0000000000000000      0 FUNC      GLOBAL DEFAULT  UND __libc_start_main@GLIBC_2.2.5 (3)
65: 0000000000000000      0 NOTYPE   WEAK     DEFAULT  UND __gmon_start__
```



Find the Entry Point

The entry point is `0x67d0`.

ELF Header:

```
Magic:  7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00 00
Class:                                     ELF64
Data:                                       2's complement, little endian
Version:                                   1 (current)
OS/ABI:                                    UNIX - System V
ABI Version:                              0
Type:                                       DYN (Shared object file)
Machine:                                  Advanced Micro Devices X86-64
Version:                                  0x1
Entry point address:                       0x67d0
Start of program headers:                  64 (bytes into file)
Start of section headers:                 140224 (bytes into file)
Flags:                                     0x0
Size of this header:                       64 (bytes)
Size of program headers:                   56 (bytes)
Number of program headers:                 13
Size of section headers:                  64 (bytes)
Number of section headers:                 30
Section header string table index:         29
```



Disassembling at the entry point

Let's try to disassemble ``which ls``.

The code at the entry point ends with a `call` to `[rip + 0x1be1e]`

This is an `RIP`-indexed instruction.

A few were possible in the 32-bit world, but this is *very common* in the 64-bit world.

Where does it go???

```
67d0: f3 0f 1e fa
67d4: 31 ed
67d6: 49 89 d1
67d9: 5e
67da: 48 89 e2
67dd: 48 83 e4 f0
67e1: 50
67e2: 54
67e3: 4c 8d 05 66 0d 01 00
67ea: 48 8d 0d ef 0c 01 00
67f1: 48 8d 3d f8 e5 ff ff
67f8: ff 15 d2 c7 01 00
67fe: f4
```

```
endbr64
xor    ebp,ebp
mov    r9,rdx
pop    rsi
mov    rdx,rsq
and    rsp,0xfffffffffffffff0
push   rax
push   rsp
lea    r8,[rip+0x10d66]
lea    rcx,[rip+0x10cef]
lea    rdi,[rip+0xffffffffffffe5f8]
call   QWORD PTR [rip+0x1c7d2]
hlt
```



Disassembling at the entry point

Take the start address of the next instruction... **0x67fe**

```
67d0: f3 0f 1e fa
67d4: 31 ed
67d6: 49 89 d1
67d9: 5e
67da: 48 89 e2
67dd: 48 83 e4 f0
67e1: 50
67e2: 54
67e3: 4c 8d 05 66 0d 01 00
67ea: 48 8d 0d ef 0c 01 00
67f1: 48 8d 3d f8 e5 ff ff
67f8: ff 15 d2 c7 01 00
67fe: f4
```

```
endbr64
xor    ebp,ebp
mov    r9,rdx
pop    rsi
mov    rdx,rsp
and    rsp,0xfffffffffffffff0
push   rax
push   rsp
lea    r8,[rip+0x10d66]
lea    rcx,[rip+0x10cef]
lea    rdi,[rip+0xffffffffffffe5f8]
call   QWORD PTR [rip+0x1c7d2]
hlt
```

Disassembling at the entry point

Take the start address of the next instruction... `0x67fe`

...and add the displacement.

```
0x1c7d2
+ 0x 67fe
-----
0x22fd0
```

This is the target of the call... but it is outside the program!

```
67d0: f3 0f 1e fa
67d4: 31 ed
67d6: 49 89 d1
67d9: 5e
67da: 48 89 e2
67dd: 48 83 e4 f0
67e1: 50
67e2: 54
67e3: 4c 8d 05 66 0d 01 00
67ea: 48 8d 0d ef 0c 01 00
67f1: 48 8d 3d f8 e5 ff ff
67f8: ff 15 d2 c7 01 00
67fe: f4
```

```
endbr64
xor     ebp,ebp
mov     r9,rdx
pop     rsi
mov     rdx,rsq
and     rsp,0xfffffffffffffff0
push    rax
push    rsp
lea     r8,[rip+0x10d66]
lea     rcx,[rip+0x10cef]
lea     rdi,[rip+0xffffffffffffe5f8]
call    QWORD PTR [rip+0x1c7d2]
hlt
```



Starting the C runtime

This call is to `__libc_start_main`.

```
int __libc_start_main(  
    int *(main) (int, char **, char **),    // This is the main function.  
    int argc,                               // Number of command line  
    arguments.  
    char ** ubp_av,                         // The command line arguments  
    (unbounded).  
    void (*init) (void),                   // The initialization function.  
    void (*fini) (void),                   // The finalization function.  
    void (*rtld_fini) (void),              // Finalize dynamic shared objects.  
    void (* stack_end));                  // The end of the stack.
```

Let's apply the calling convention.



Starting the C runtime

This call is to `__libc_start_main`.

```
int __libc_start_main(  
    int *(main) (int, char **, char **),    // rdi  
    int argc,                                // rsi  
    char ** ubp_av,                        // rdx  
    void (*init) (void),                    // rcx  
    void (*fini) (void),                   // r8  
    void (*rtld_fini) (void),              // r9  
    void (* stack_end));                   // on stack
```



Starting the C runtime

At entry (from the loader):

- `RDX` contains the address of the destructor function call handler for the dynamic linker, `_dl_fini`.
- The stack contains `argc`, `argv`, and `envp`, with `argc` on top.

When we call `main`, we at least want:

- `EDI` to contain `argc`
- `ESI` to contain the pointer to `argv`
- `EDX` to contain the pointer to `envp`

```
67d0: endbr64
67d4: xor     ebp,ebp
67d6: mov     r9,rdx
67d9: pop     rsi
67da: mov     rdx,rsp
67dd: and     rsp,0xfffffffffffffff0
67e1: push    rax
67e2: push    rsp
67e3: lea     r8,[rip+0x10d66]
67ea: lea     rcx,[rip+0x10cef]
67f1: lea     rdi,[rip+0xffffffffffffe5f8]
67f8: call    QWORD PTR [rip+0x1c7d2]
67fe: hlt
```



Starting the C runtime

```
int __libc_start_main(  
    int *(main) (int, char **, char **),    // rdi  
    int argc,  
                                           // rsi  
    char ** ubp_av,  
                                           // rdx  
    void (*init) (void),  
        // rcx  
    void (*fini) (void),  
        // r8  
    void (*rtld_fini) (void),  
        // r9  
    void (* stack_end));  
    // on stack
```


Let's map assembly instructions to arguments.

```
67d0: endbr64  
67d4: xor     ebp,ebp  
67d6: mov     r9,rdx  
67d9: pop     rsi  
67da: mov     rdx,rsp  
67dd: and     rsp,0xfffffffffffffff0  
67e1: push    rax  
67e2: push    rsp  
67e3: lea     r8,[rip+0x10d66]  
67ea: lea     rcx,[rip+0x10cef]  
67f1: lea     rdi,[rip+0xffffffffffffe5f8]  
67f8: call    QWORD PTR [rip+0x1c7d2]  
67fe: hlt
```

Starting the C runtime

```
int __libc_start_main(
    int *(main) (int, char **, char **),    // rdi
    int argc,
                                           // rsi
    char ** ubp_av,
                                           // rdx
    void (*init) (void),
        // rcx
    void (*fini) (void),
        // r8
    void (*rtld_fini) (void),
        // rcx
    void (* stack_end));
    // on stack
```

This is a branch protection instruction. Think of it as a no-op. Technically it is!



```
67d0: endbr64
67d4: xor     ebp,ebp
67d6: mov     r9,rdx
67d9: pop     rsi
67da: mov     rdx,rsp
67dd: and     rsp,0xfffffffffffffff0
67e1: push    rax
67e2: push    rsp
67e3: lea     r8,[rip+0x10d66]
67ea: lea     rcx,[rip+0x10cef]
67f1: lea     rdi,[rip+0xffffffffffffe5f8]
67f8: call    QWORD PTR [rip+0x1c7d2]
67fe: hlt
```

Starting the C runtime

```
int __libc_start_main(  
    int *(main) (int, char **, char **),    // rdi  
    int argc,  
                                             // rsi  
    char ** ubp_av,  
                                             // rdx  
    void (*init) (void),  
        // rcx  
    void (*fini) (void),  
        // r8  
    void (*rtld_fini) (void),  
        // r9  
    void (* stack_end));  
    // on stack
```

```
67d0: endbr64  
67d4: xor     ebp,ebp  
67d6: mov     r9,rdx  
67d9: pop     rsi  
67da: mov     rdx,rsi  
67dd: and     rsp,0xfffffffffffffff0  
67e1: push    rax  
67e2: push    rsp  
67e3: lea     r8,[rip+0x10d66]  
67ea: lea     rcx,[rip+0x10cef]  
67f1: lea     rdi,[rip+0xffffffffffffe5f8]  
67f8: call    QWORD PTR [rip+0x1c7d2]  
67fe: hlt
```

Starting the C runtime

```
int __libc_start_main(  
    int *(main) (int, char **, char **),    // rdi  
    int argc,                               // rsi  
    char ** ubp_av,                          // rdx  
    void (*init) (void),                    // rcx  
    void (*fini) (void),                    // r8  
    void (*rtld_fini) (void),               // r9  
    void (* stack_end));  
    // on stack
```

```
67d0: endbr64  
67d4: xor     ebp,ebp  
67d6: mov     r9,rdx  
67d9: pop     rsi  
67da: mov     rdx,rsi  
67dd: and     rsp,0xfffffffffffffff0  
67e1: push    rax  
67e2: push    rsp  
67e3: lea     r8,[rip+0x10d66]  
67ea: lea     rcx,[rip+0x10cef]  
67f1: lea     rdi,[rip+0xffffffffffffe5f8]  
67f8: call    QWORD PTR [rip+0x1c7d2]  
67fe: hlt
```

Starting the C runtime

```
int __libc_start_main(  
    int *(main) (int, char **, char **),    // rdi  
    int argc,                               // rsi  
    char ** ubp_av,                          // rdx  
    void (*init) (void),                    // rcx  
    void (*fini) (void),                    // r8  
    void (*rtld_fini) (void),               // r9  
    void (* stack_end));  
    // on stack
```

```
67d0: endbr64  
67d4: xor     ebp,ebp  
67d6: mov     r9,rdx  
67d9: pop     rsi  
67da: mov     rdx,rsp  
67dd: and     rsp,0xfffffffffffffff0  
67e1: push    rax  
67e2: push    rsp  
67e3: lea     r8,[rip+0x10d66]  
67ea: lea     rcx,[rip+0x10cef]  
67f1: lea     rdi,[rip+0xffffffffffffe5f8]  
67f8: call    QWORD PTR [rip+0x1c7d2]  
67fe: hlt
```

Starting the C runtime

```
int __libc_start_main(
    int *(main) (int, char **, char **),    // rdi
    int argc,
                                           // rsi
    char ** ubp_av,
                                           // rdx
    void (*init) (void),
        // rcx
    void (*fini) (void),
        // r8
    void (*rtld_fini) (void),
        // r9
    void (* stack_end));
    // on stack
```

```
67d0: endbr64
67d4: xor     ebp,ebp
67d6: mov     r9,rdx
67d9: pop     rsi
67da: mov     rdx,rsp
67dd: and     rsp,0xfffffffffffffff0
67e1: push    rax
67e2: push    rsp
67e3: lea     r8,[rip+0x10d66]
67ea: lea     rcx,[rip+0x10cef]
67f1: lea     rdi,[rip+0xfffffffffffffe5f8]
67f8: call    QWORD PTR [rip+0x1c7d2]
67fe: hlt
```

16-bit stack
alignment

Starting the C runtime

```
int __libc_start_main(
    int *(main) (int, char **, char **),    // rdi
    int argc,                                // rsi
    char ** ubp_av,                          // rdx
    void (*init) (void),                    // rcx
    void (*fini) (void),                    // r8
    void (*rtld_fini) (void),               // r9
    void (* stack_end));
    // on stack
```

```
67d0: endbr64
67d4: xor     ebp,ebp
67d6: mov     r9,rdx
67d9: pop     rsi
67da: mov     rdx,rsp
67dd: and     rsp,0xfffffffffffffffff0
67e1: push    rax
67e2: push    rsp
67e3: lea     r8,[rip+0x10d66]
67ea: lea     rcx,[rip+0x10cef]
67f1: lea     rdi,[rip+0xffffffffffffe5f8]
67f8: call    QWORD PTR [rip+0x1c7d2]
67fe: hlt
```

This messes
it up!

Starting the C runtime

```
int __libc_start_main(
    int *(main) (int, char **, char **),    // rdi
    int argc,                                // rsi
    char ** ubp_av,                          // rdx
    void (*init) (void),                    // rcx
    void (*fini) (void),                    // r8
    void (*stack_end) (void),               // r9
    void (* stack_end));
    // on stack
```

Now the stack end (RSP) is at the top of an aligned stack.

```
67d0: endbr64
67d4: xor     ebp,ebp
67d6: mov     r9,rdx
67d9: pop     rsi
67da: mov     rdx,rsp
67dd: and     rsp,0xfffffffffffffff0
67e1: push    rax
67e2: push    rsp
67e3: lea     r8,[rip+0x10d66]
67ea: lea     rcx,[rip+0x10cef]
67f1: lea     rdi,[rip+0xfffffffffffffe5f8]
67f8: call    QWORD PTR [rip+0x1c7d2]
67fe: hlt
```

This fixes it!

Starting the C runtime

```
int __libc_start_main(  
    int *(main) (int, char **, char **),    // rdi  
    int argc,  
                                           // rsi  
    char ** ubp_av,  
                                           // rdx  
    void (*init) (void),  
        // rcx  
    void (*fini) (void),  
        // r8  
    void (*rtld_fini) (void),  
        // r9  
    void (* stack_end));  
    // on stack
```

```
67d0: endbr64  
67d4: xor     ebp,ebp  
67d6: mov     r9,rdx  
67d9: pop     rsi  
67da: mov     rdx,rsi  
67dd: and     rsp,0xfffffffffffffff0  
67e1: push    rax  
67e2: push    rsp  
67e3: lea     r8,[rip+0x10d66]  
67ea: lea     rcx,[rip+0x10cef]  
67f1: lea     rdi,[rip+0xffffffffffffe5f8]  
67f8: call    QWORD PTR [rip+0x1c7d2]  
67fe: hlt
```

Starting the C runtime

```
int __libc_start_main(  
    int *(main) (int, char **, char **),    // rdi  
    int argc,  
                                           // rsi  
    char ** ubp_av,  
    void (*init) (void),                    // rdx  
                                           // rcx  
    void (*fini) (void),  
                                           // r8  
    void (*rtld_fini) (void),  
                                           // r9  
    void (* stack_end));  
    // on stack
```

```
67d0: endbr64  
67d4: xor     ebp,ebp  
67d6: mov     r9,rdx  
67d9: pop     rsi  
67da: mov     rdx,rsi  
67dd: and     rsp,0xfffffffffffffff0  
67e1: push    rax  
67e2: push    rsp  
67e3: lea     r8,[rip+0x10d66]  
67ea: lea     rcx,[rip+0x10cef]  
67f1: lea     rdi,[rip+0xffffffffffffe5f8]  
67f8: call    QWORD PTR [rip+0x1c7d2]  
67fe: hlt
```

Starting the C runtime

```
int __libc_start_main(  
    int *(main) (int, char **, char **),    // rdi  
    int argc,  
    char ** ubp_av,                          // rsi  
    void (*init) (void),                     // rdx  
    void (*fini) (void),                     // rcx  
    void (*rtld_fini) (void),                // r8  
    void (*stack_end));                     // r9  
    // on stack
```

```
67d0: endbr64  
67d4: xor     ebp,ebp  
67d6: mov     r9,rdx  
67d9: pop     rsi  
67da: mov     rdx,rsi  
67dd: and     rsp,0xfffffffffffffff0  
67e1: push    rax  
67e2: push    rsp  
67e3: lea     r8,[rip+0x10d66]  
67ea: lea     rcx,[rip+0x10cef]  
67f1: lea     rdi,[rip+0xffffffffffffe5f8]  
67f8: call    QWORD PTR [rip+0x1c7d2]  
67fe: hlt
```

Starting the C runtime

```
int __libc_start_main(  
→ int *(main) (int, char **, char **),           // rdi  
  int argc,  
                                     // rsi  
  char ** ubp_av,  
                                     // rdx  
  void (*init) (void),  
    // rcx  
  void (*fini) (void),  
    // r8  
  void (*rtld_fini) (void),  
    // r9  
  void (* stack_end));  
    // on stack
```

```
67d0: endbr64  
67d4: xor    ebp,ebp  
67d6: mov    r9,rdx  
67d9: pop    rsi  
67da: mov    rdx,rsp  
67dd: and    rsp,0xfffffffffffffff0  
67e1: push   rax  
67e2: push   rsp  
67e3: lea    r8,[rip+0x10d66]  
67ea: lea    rcx,[rip+0x10cef]  
→ 67f1: lea    rdi,[rip+0xfffffffffffffe5f8]  
67f8: call   QWORD PTR [rip+0x1c7d2]  
67fe: hlt
```

Starting the C runtime

```
int __libc_start_main(  
→ int *(main) (int, char **, char **),      // rdi  
  int argc,                                // rsi  
  char ** ubp_av,                          // rdx  
  void (*init) (void),                    // rcx  
  void (*fini) (void),                    // r8  
  2's comp. 0x10000,                      // r9  
  void (*rtld_fini) (void),               - 0x  
  e5f8,                                    // r9  
  void (* stack_end));                    0x  
  1a08 (disp) // on stack
```

```
67d0: endbr64  
67d4: xor    ebp,ebp  
67d6: mov    r9,rdx  
67d9: pop    rsi  
67da: mov    rdx,rsp  
67dd: and    rsp,0xfffffffffffffff0  
67e1: push   rax  
67e2: push   rsp  
67e3: lea    r8,[rip+0x10d66]  
67ea: lea    rcx,[rip+0x10cef]  
→ 67f1: lea    rdi,[rip+0xffffffffffffe5f8]  
67f8: call   QWORD PTR [rip+0x1c7d2]  
67fe: hlt
```

Starting the C runtime

```
int __libc_start_main(  
→ int *(main) (int, char **, char **),      // rdi  
  int argc,  
                                     // rsi  
  char ** ubp_av,  
                                     // rdx  
  void (*init) (void),  
    // rcx  
  void (*fini) (void),  
    // r8  
main:  
  void (*rtld_fini) (void),  
0x67f8 (rip) // r9  
  void (*stack_end));  
0x1a08 (disp) // on stack  
0x4df0
```

```
67d0: endbr64  
67d4: xor    ebp,ebp  
67d6: mov    r9,rdx  
67d9: pop    rsi  
67da: mov    rdx,rsi  
67dd: and    rsp,0xfffffffffffffff0  
67e1: push   rax  
67e2: push   rsp  
67e3: lea    r8,[rip+0x10d66]  
67ea: lea    rcx,[rip+0x10cef]  
→ 67f1: lea    rdi,[rip+0xfffffffffffffe5f8]  
67f8: call   QWORD PTR [rip+0x1c7d2]  
67fe: hlt
```


Starting the C runtime

```
int __libc_start_main(  
→ int *(main) (int, char **, char **),           // rdi  
  int argc,  
                                     // rsi  
  char ** ubp_av,  
                                     // rdx  
  void (*init) (void),  
    // rcx  
  void (*fini) (void),  
    // r8  
main:                                0x67f8 (rip)  
  void (*rtld_fini) (void),           - 0x1a08 (disp)  
    // r9  
  void (* stack_end));              0x4df0  
                                     // on stack  
>>> hex(0x67f8 - (0x10000-0xe5f8))  
'0x4df0'
```

```
67d0: endbr64  
67d4: xor    ebp,ebp  
67d6: mov    r9,rdx  
67d9: pop    rsi  
67da: mov    rdx,rsp  
67dd: and    rsp,0xfffffffffffffff0  
67e1: push   rax  
67e2: push   rsp  
67e3: lea    r8,[rip+0x10d66]  
67ea: lea    rcx,[rip+0x10cef]  
→ 67f1: lea    rdi,[rip+0xfffffffffffffe5f8]  
67f8: call   QWORD PTR [rip+0x1c7d2]  
67fe: hlt
```

Starting the C runtime

```
➡ int __libc_start_main(  
    int *(main) (int, char **, char **),    // rdi  
    int argc,  
                                           // rsi  
    char ** ubp_av,  
                                           // rdx  
    void (*init) (void),  
        // rcx  
    void (*fini) (void),  
        // r8  
    void (*rtld_fini) (void),  
        // r9  
    void (* stack_end));  
    // on stack
```

```
67d0: endbr64  
67d4: xor    ebp,ebp  
67d6: mov    r9,rdx  
67d9: pop    rsi  
67da: mov    rdx,rsp  
67dd: and    rsp,0xfffffffffffffffff0  
67e1: push   rax  
67e2: push   rsp  
67e3: lea    r8,[rip+0x10d66]  
67ea: lea    rcx,[rip+0x10cef]  
67f1: lea    rdi,[rip+0xffffffffffffe5f8]  
➡ 67f8: call   QWORD PTR [rip+0x1c7d2]  
67fe: hlt
```

Starting the C runtime

```
int __libc_start_main(
    int *(main) (int, char **, char **),    // rdi
    int argc,
                                           // rsi
    char ** ubp_av,
                                           // rdx
    void (*init) (void),
        // rcx
    void (*fini) (void),
        // r8
    void (*rtld_fini) (void),
    void (*stack_end));
    // on stack
```

This does not return! Remember that the top-level process (`_start`) needs to call `sys_exit`.

```
67d0: endbr64
67d4: xor     ebp,ebp
67d6: mov     r9,rdx
67d9: pop     rsi
67da: mov     rdx,rsp
67dd: and     rsp,0xfffffffffffffffff0
67e1: push    rax
67e2: push    rsp
67e3: lea     r8,[rip+0x10d66]
67ea: lea     rcx,[rip+0x10cef]
67f1: lea     rdi,[rip+0xfffffffffffffe5f8]
67f8: call    QWORD PTR [rip+0x1c7d2]
67fe: hlt
```





Starting the C runtime

Not always so indirect.

5cff70:	f3 0f 1e fa	endbr64
5cff74:	31 ed	xor ebp,ebp
5cff76:	49 89 d1	mov r9,rdx
5cff79:	5e	pop rsi
5cff7a:	48 89 e2	mov rdx,rsi
5cff7d:	48 83 e4 f0	mov rdx,rsi
5cff81:	50	and rsp,0xfffffffffffffffff0
5cff82:	54	push rax
5cff83:	49 c7 c0 00 95 67 00	push rsp
5cff8a:	48 c7 c1 90 94 67 00	mov r8,0x679500
→ 5cff91:	48 c7 c7 60 f9 4c 00	mov rcx,0x679490
5cff98:	ff 15 5a 90 26 00	mov rdi,0x4cf960
5cff9e:	f4	call QWORD PTR [rip+0x26905a]
		hlt

A Structuring Example



XOn / XOff Flow Control

RS232 (character oriented stream) flow control.

Essentially: Who gets to talk.

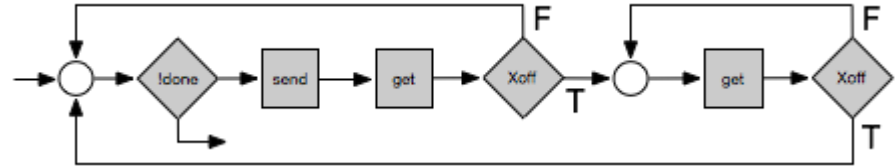
When the receiver buffer fills to the point it cannot accept any more data it sends an XOff (Transmit Off) to the transmitter. When the transmitter sees the XOff character, it stops transmitting. When the receiver can again accept data, it sends XOn. When the transmitter sees XOn, it resumes sending.

Normally we have:

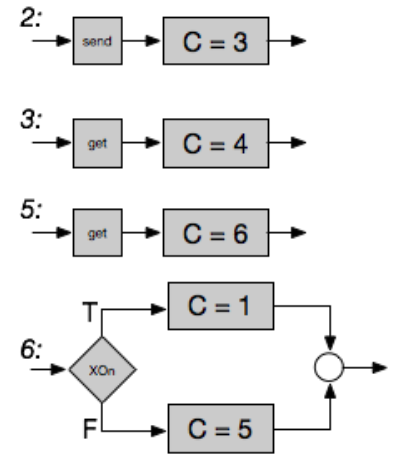
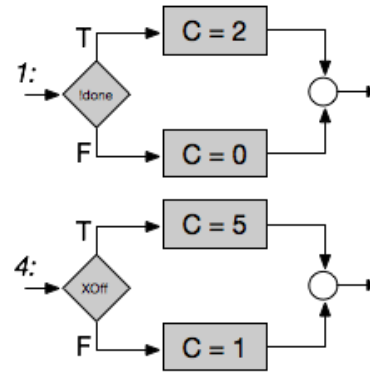
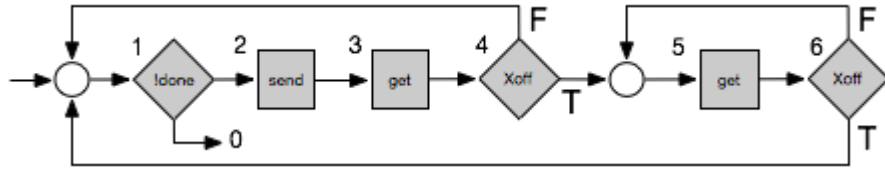
- XOn = CTRL+Q = 0x11
- XOff = CTRL+S = 0x13

XOn / XOff Flow Control

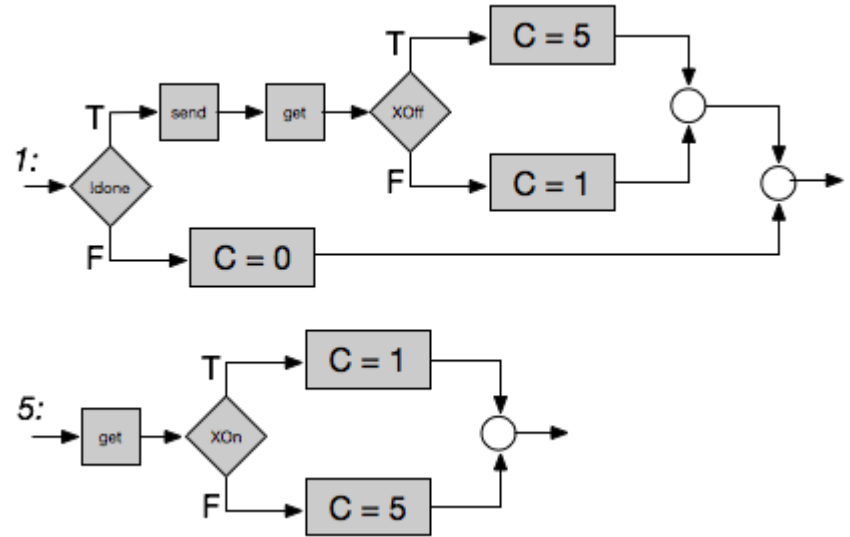
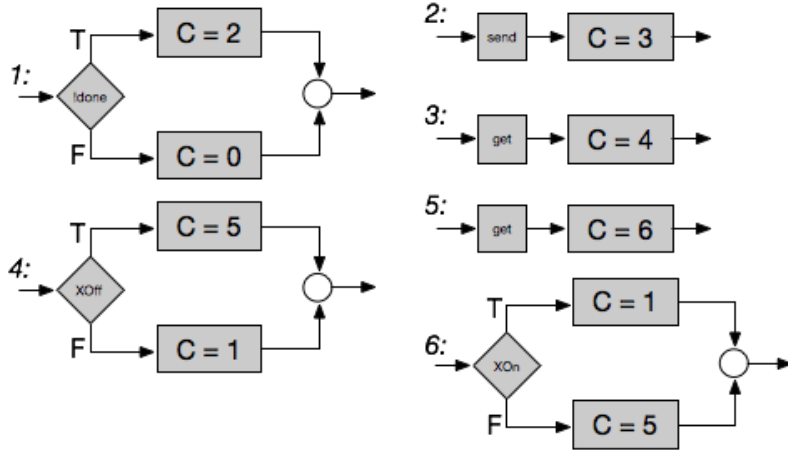
```
for (i = 0; i < buflen; ++i) {  
    transmit_byte();  
    b = receive_byte();  
    if (b != XOFF) continue;  
    while (b != XON) {  
        b = receive_byte();  
    }  
}
```



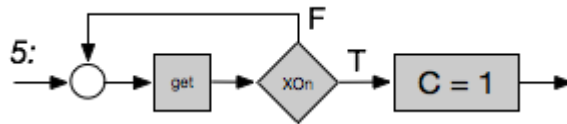
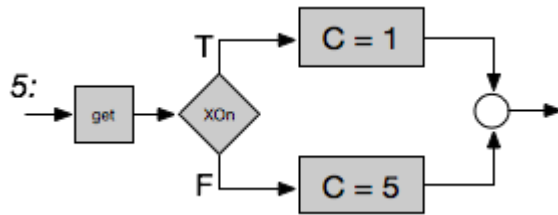
XOn / XOff Flow Control



XOn / XOff Flow Control

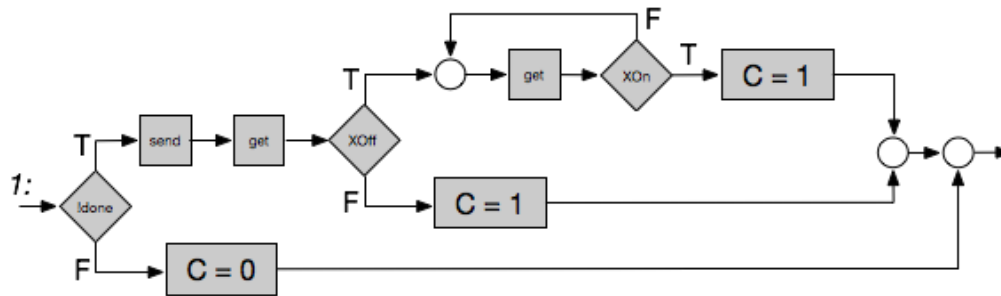
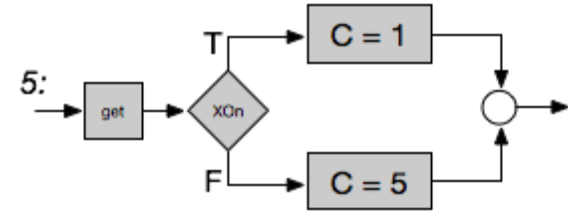
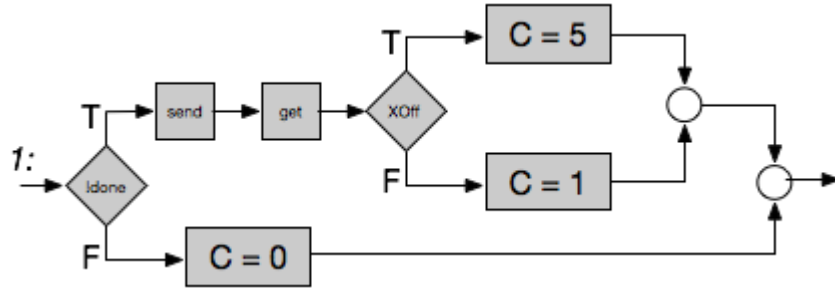


XOn / XOff Flow Control

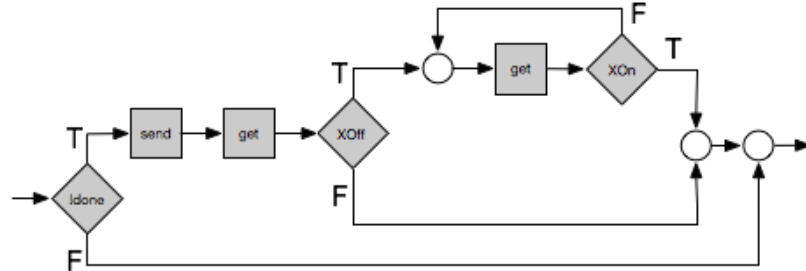
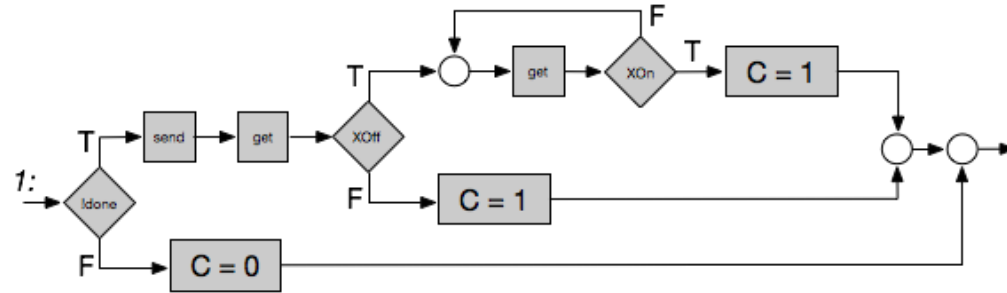


```
do {  
    b = receive_byte();  
} while (b != XON);
```

XOn / XOff Flow Control



XOn / XOff Flow Control



```

for (i = 0; i < buflen; ++i) {
    transmit_byte();
    b = receive_byte();
    if (b == XOFF) {
        while (b != XON) {
            b = receive_byte();
        }
    }
}
  
```

**Next time:
Control Flow**