



# **CSC 6580**

# **Spring 2020**

Instructor: Stacy Prowell

# Decidability (Again)

---



# Hailstone Numbers

Given integer  $n_0 > 0$ , let a sequence of integers be defined as follows.

$$\begin{cases} n_{i+1} = 3n_i + 1 & \text{if } n_i \text{ is odd} \\ n_{i+1} = n_i / 2 & \text{if } n_i \text{ is even} \end{cases}$$

This sequence terminates when  $n_i = 1$ .

**The Collatz Conjecture:** This sequence terminates for any positive integer  $n_0$ .

The number sequence is known as the "hailstone numbers."

```
def hailstone(start: int):  
    print(start, end='')  
    while start != 1:  
        print(', ', end='')  
        if start & 1 == 1:  
            start = start * 3 + 1  
        else:  
            start //= 2  
        print(start, end='')
```

```
>>> hailstone(23)  
23, 70, 35, 106, 53, 160, 80, 40, 20, 10, 5,  
16, 8, 4, 2, 1
```



# Hailstone Numbers

The C program at right *terminates* iff the input is a terminating hailstone number. Otherwise it does *not terminate*.

It is *not known* if the Collatz conjecture is true. It is an open question.

Is termination of this program *decidable*?

```
#include <stdint.h>

void hailstone(uint64_t start) {
    while (start != 1) {
        if (start & 1) {
            start = start * 3 + 1;
        } else {
            start = start / 2;
        }
    }
}
```



# Hailstone Numbers

Here is a test harness for the function.

*Estimated time for various limits, assuming all values terminate!*

All `uint16_t`: instant

All `uint32_t`: 47 minutes

All `uint64_t`: 393,000 years

But what if some do not terminate? How would we know? Keep waiting?

```
#include <stdint.h>

#define LIMIT 0xffff

extern void hailstone(uint64_t);

int main(int argc, char * argv[]) {
    for (uint64_t start = 1;
        start < LIMIT;
        ++start) {
        hailstone(start);
    }
}
```



# Hailstone Numbers

It's decidable.

```
#include <stdint.h>

#define LIMIT 0xffff

extern void hailstone(uint64_t);

int main(int argc, char * argv[]) {
    for (uint64_t start = 1;
        start < LIMIT;
        ++start) {
        hailstone(start);
    }
}
```



# Hailstone Numbers

It's decidable.

This is basically a (very large) finite state machine.

In any sufficiently long sequence we will either (1) reach one and terminate, or (2) repeat a number and we will know this does not terminate. So with sufficient memory (a lot!) and sufficient time (a lot!) we can always decide.

```
#include <stdint.h>

#define LIMIT 0xffff

extern void hailstone(uint64_t);

int main(int argc, char * argv[]) {
    for (uint64_t start = 1;
        start < LIMIT;
        ++start) {
        hailstone(start);
    }
}
```



# Why is this relevant?

Most things about most programs are decidable.

It is *hard* to write a concrete program whose termination is undecidable, and not just hard to figure out or unknown. Give it a try!

Even in the extreme example of the Collatz Conjecture, whose answer is *unknown* we are still decidable, even if the general Collatz Conjecture is *unprovable* or even *provably undecidable*.



# A (Ridiculously) Short Introduction to the Very Important Technique of Program Slicing\*

---

\* Abridged version



# A Program

Here's a short program.

Under what conditions does it terminate?

We can start to figure that out by constructing the *control slice*.

We pick a point in the program and a value (or values) we want, and then discard anything that does not affect those values at that point.

```
start:
    xor    eax, eax
    mov    rbx, START
.loop:
    add    rbx, STRIDE
    inc    rax
    cmp    rbx, END
    jne    .loop
    ret
```



# A Program

If this sounds familiar, it has a lot in common with live variable analysis.

We want to know *when this program terminates*.

It terminates when we reach the return.

We reach the return when the branch is not taken.

```
start:
    xor    eax, eax
    mov    rbx, START
.loop:
    add    rbx, STRIDE
    inc    rax
    cmp    rbx, END
    jne    .loop           ; ZF
must be set
ret
```



# A Program

Need to know when **RBX** is equal to **END**, so we only care about **RBX**. We select the statements that affect **RBX**.

This is like live variable analysis with only **RAX** live.

```
start:
    xor     eax, eax
    mov     rbx, START           ; needed
.loop:
    add     rbx, STRIDE          ; needed
    inc     rax
    cmp     rbx, END             ; rbx must be END
    jne     .loop                ; ZF
must be set
    ret
```



# A Program

These are the instructions that affect the program's control flow (which is all we care about). This is called the **control slice**.

```
    mov    rbx, START
.loop:
    add    rbx, STRIDE
    cmp    rbx, END
    jne    .loop
```



# A Program

Let's simplify the variables.

We need to know when  $RBX = E$ .

Initially  $RBX = T$ . Assume the loop runs  $n$  times, then we add  $nS$  to  $RBX$ , but we do this for a 64-bit register. So we have:

$$T + nS \cong E \pmod{2^{64}}$$

```
mov    rbx, T
.loop:
add     rbx, S
cmp     rbx, E
jne     .loop
```



# A Program

$$T + nS = E \pmod{2^{64}}$$

We know  $T$ ,  $S$ , and  $E$ . If we can find an  $n$  that makes the above true, then the loop halts. If we cannot, then the loop doesn't halt.

We have the following.

$$2^{64} \mid T + nS - E$$

```
mov    rbx, T
.loop:
add     rbx, S
cmp     rbx, E
jne     .loop
```



# A Program

$$2^{64} \mid T + nS - E$$

Is there an  $n$  that *satisfies* this expression? Better yet, is there a program that can find it for us?

```
mov    rbx, T
.loop:
add     rbx, S
cmp     rbx, E
jne     .loop
```



# Satisfiability

---



# Satisfiability

Given a *formula* is there some *model* (assignment of values to variables) that makes the formula *true*? If so, the formula is **satisfiable**. Otherwise it is not. (If all models make the formula true, the formula is **valid**.)

In propositional logic (variables, truth constants, and logical connectives but no quantifiers) satisfiability is *decidable*.

In first-order (predicate) logic (quantifiers) satisfiability is *not decidable*.

$$\neg p \vee p$$

valid

$$p \wedge \neg p$$

invalid / not  
satisfiable

$$p \implies q$$

satisfiable



# Satisfiability

Boolean satisfiability (SAT) is the *first* problem proven to be NP-complete.

Check it by plugging in the answer. But finding the answer might be really hard...



# SAT Solvers

Convert a Boolean expression to conjunctive normal form (CNF).

$$\neg p \vee (\neg q \wedge r) \quad \Longrightarrow \quad (\neg p \vee \neg q) \wedge (\neg p \vee r)$$

Now try to find an assignment that makes all the individual clauses true.

There can be multiple solutions.

Which one a SAT solver chooses can be *random*, and vary from run to run.  
Some allow specifying a random seed so the result is repeatable.

$$\begin{aligned} p = \text{T} \wedge q = \text{F} \wedge r = \text{T} \\ p = \text{F} \wedge q = \text{F} \wedge r = \text{F} \\ p = \text{F} \wedge q = \text{T} \wedge r = \text{F} \\ \vdots \end{aligned}$$

**Remember:** They give you *a* solution, not necessarily *the* solution.



# SAT Solvers

Lots of them. Some well-known ones are:

- MiniSat <http://minisat.se/>
- Glucose <https://www.labri.fr/perso/lrsimon/glucose/>
- Lingeling <http://fmv.jku.at/lingeling/>

Also: PySAT <https://pysathq.github.io/> provides a Python interface to a lot of different SAT solvers... the above three included.

# Satisfiability Modulo Theories

---



# Integers, Bit Arrays, and Other Stuff

Boolean satisfiability is find and all, but how do we use this to prove things about programs?

Well, start with a Boolean satisfiability problem.

$$(\neg p \vee \neg q) \wedge (\neg p \vee r)$$

Now replace some of the variables with expressions about integers.

$$(\neg(x \geq 15)) \vee \neg q \wedge (\neg(x \geq 15)) \vee r)$$

These additions are called *background theories* and we can think of the result as a *tiered* satisfiability problem. Solve the outer problem, then this constrains the inner problems: *satisfiability modulo theories*.



# Theories

Some theories are decidable.

- Pressburger Arithmetic

Some are not.

- Peano Arithmetic

Many are very useful for CS.

- Arrays, bit vectors, lists, ...



**Z3**

—



# Z3

```
$ pip3 install z3-solver
```

Z3 is an open source SMT solver developed by Microsoft.

It has a nice Python interface.

<https://ericpony.github.io/z3py-tutorial/guide-examples.htm>



# Z3

Let's try a simple problem in Z3.

$$\neg p \vee (\neg q \wedge r)$$

```
from z3 import *  
p = Bool('p')  
q = Bool('q')  
r = Bool('r')  
s = Solver()  
s.add(Or(Not(p), And(Not(q), r)))  
print(s.check())  
print(s.model())
```



# Z3

Let's try a simple problem in Z3.

$$\neg p \vee (\neg q \wedge r)$$

[p = False, q = False, r = True]

```
from z3 import *  
p = Bool('p')  
q = Bool('q')  
r = Bool('r')  
s = Solver()  
s.add(Or(Not(p), And(Not(q), r)))  
print(s.check())  
print(s.model())
```



# SMT-LIB

There is a standard language for SMT solvers, called SMT-LIB.

Work with Z3 using SMT-LIB online:

<https://rise4fun.com/Z3>

Find a tutorial guide here:

<https://rise4fun.com/z3/tutorial>

```
(declare-const p Bool)
(declare-const q Bool)
(declare-const r Bool)
(assert (or (not p) (and (not q) r)))
(check-sat)
(get-model)
```

# Solving Termination with Z3

---



# Solving

Here is an SMT-LIB program in Z3 for solving our earlier termination problem.

Recall that we need an  $n$  such that:

$$2^{64} \mid T + nS - E$$

Trying this when  $T = 7$ ,  $S = 4$ , and  $E = 22$ ... we find the model is *unsatisfiable*. The program runs forever.

```
(declare-fun n () Int)
(declare-fun T () Int)
(declare-fun S () Int)
(declare-fun E () Int)

; Our constraints for a solution.
(assert (= T 7))
(assert (= S 4))
(assert (= E 22))
(assert (> n 0))
(assert
  (= 0
    (mod
      (+ T (* n S) (- E)) (^ 2 64))))

; Solve
(check-sat)
(get-model)
```



# Solving

Trying this again, but changing  $E = 23$ ... we find the model is now *satisfiable*. But (when writing these slides) the value for  $n$  shows:

4,611,686,018,427,387,908

That's... awesome and all, but what about 4?

$$7 + 4 * 4 = 23$$

Hey, it found *a* solution. You need to add other constraints if you want the *first* value.

```
from z3 import *
n = Int('n')
T = Int('T')
S = Int('S')
E = Int('E')
s = Solver()
s.add(T == 7)
s.add(S == 4)
s.add(E == 23)
s.add(n > 0)
s.add(0 == ((T + n*S - E) % 2**64))
print(s.check())
print(s.model())
```





# TERMINATOR

Byron Cook created a tool called TERMINATOR to prove termination of programs.

We can “reduce” lots of problems to termination. For instance, if the correct password is entered, the program terminates; if a particular string is printed, the program terminates; etc.

TERMINATOR tried to find a well-founded relation on variables of interest to prove termination--just like we did earlier.

There is a YouTube interview with him about this:

<https://channel9.msdn.com/Shows/Going+Deep/Byron-Cook-Inside-Terminator>

This is now part of T2: <http://mmjb.github.io/T2/>

# Concolic Execution

---



# Concrete Execution

How long will it take to find the error in testing if we just generate random inputs? This is *concrete execution*.

```
void f(int x, int y) {  
    int z = 2*y;  
    if (x == 100000) {  
        if (x < z) {  
            assert(0); /* error */  
        }  
    }  
}
```



# Symbolic Execution

We can *symbolically execute* the program to find the error.

```
void f(int x, int y) {  
    int z = 2*y;  
    if (x == 100000) {  
        if (x < z) {  
            assert(0); /* error */  
        }  
    }  
}
```



# Symbolic Execution

```
void f(int x, int y) {  
    int z = 2*y;  
    if (x == 100000) {  
        if (x < z) {  
            assert(0); /* error */  
        }  
    }  
}
```

```
z = 2*y  
x == 100000:  
    100000 < 2*y → error  
x != 100000:  
    pass
```

This can be very hard and may take a long time for large programs.



# Concolic Execution

1. Identify a set of variables as *inputs*.
2. Instrument the program to *trace* modifications to inputs or control flow changes.
3. Pick some initial input *values*.
4. Do:
  - a. Execute the program on the input values.
  - b. Symbolically re-execute the program *on the trace* and generate *constraints* on inputs and path conditions.
  - c. Negate the last path condition (not already negated) to explore (potentially) an alternate path.
  - d. Call a satisfiability solver on the constraints to generate a new input values. If not satisfied, go back and negate another path condition. If no more, halt.



# Instrumentation? Pin

Pin supports dynamic instrumentation of binaries on Linux, Windows, and macOS.

<https://software.intel.com/en-us/articles/pin-a-dynamic-binary-instrumentation-tool>

You *do not need* to re-compile the target program to instrument it; Pin is essentially Just-in-Time (JIT) instrumentation.



# Concolic Execution

```
void f(int x, int y) {  
    int z = 2*y;  
    if (x == 100000) {  
        if (x < z) {  
            assert(0); /* error */  
        }  
    }  
}
```

- The inputs are  $x$  and  $y$
- Instrument the code
- Let  $x = y = 1$ .
- Execute  $f(1,1)$
- The branch is not taken ( $x \neq 100000$ )





# Concolic Execution

```
void f(int x, int y) {  
    int z = 2*y;  
    if (x == 100000) {  
        if (x < z) {  
            assert(0); /* error */  
        }  
    }  
}
```

- Now symbolically execute
- Set  $z = 2*y$
- Path condition is  $x \neq 100000$



# Concolic Execution

```
void f(int x, int y) {  
    int z = 2*y;  
    if (x == 100000) {  
        if (x < z) {  
            assert(0); /* error */  
        }  
    }  
}
```

- Negate last path condition:  
`not(x != 100000)`
- Call SMT solver to get a valid `x` and `y`
- Get `x = 100000, y = 1`
- Execute `f(100000,1)`
- The outer branch is taken, but the inner branch is not



# Concolic Execution

```
void f(int x, int y) {  
    int z = 2*y;  
    if (x == 100000) {  
        if (x < z) {  
            assert(0); /* error */  
        }  
    }  
}
```

- Now symbolically execute
- Set  $z = 2*y$
- Path conditions are now  $x == 100000$  and  $x \geq z$ , which becomes  $x \geq 2*y$



# Concolic Execution

```
void f(int x, int y) {  
    int z = 2*y;  
    if (x == 100000) {  
        if (x < z) {  
            assert(0); /* error */  
        }  
    }  
}
```

- Negate last path condition:  
`not(x >= 2*y)`
- Run SMT solver to get valid `x` and `y`
- Get `x = 100000, y = 500001`
- Execute `f(100000, 500001)`
- Both the outer and inner branch are now taken



# Concolic Execution with Pin and Z3

<https://triton.quarkslab.com/>

---

**Next time:**  
**Last class!**  
**Discussion of Final Exam**